

Data Structures III: Tibbles



UNIVERSITY OF
SAN FRANCISCO

Abbie M Popa

BSDS 100 - Intro to Data Science with R



- Tibbles
- Tibbles versus Data Frames



- Dataframes can be fickle, with their `stringsAsFactors = FALSE` requirements, etc.
- Tibbles - from the `tibble` package - **are data frames**, with some minor modifications which make operating with tibbles a more pleasant and worry-free experience than dealing with raw dataframes.
- Creating a tibble is the same as creating a dataframe, but instead we use the function `tibble()` instead of `data.frame()`.
- Coercing a dataframe (or other data structure) is done with `as_tibble()` function



- Tibbles have some convenient advantages
- When importing (e.g., `CSV`'s) or coercing data into a `tibble()`
 - 1 **will not** automatically convert strings to factors
 - 2 **will not** change the names of variables
 - 3 **will not** create row names
 - 4 **will not** do any partial matching when subsetting
 - 5 when subsetting a tibble, it is consistent in returning a tibble or a vector
- Subsetting tibbles can be done with either `$` or `[]`



- Recall that a dataframe is a list, which means that `typeof(myDataFrame)` will output a list
- Instead use `is.data.frame()`
- An object can be coerced to a (base) dataframe using `as.data.frame()` or `tibble::as_data_frame()`

Subsetting Dataframes



Dataframes (both base and tibbles) can be subset like lists or like matrices

```
> (my_data_frame <- data.frame(x = 1:3, y = -4:-6, z = LETTERS[1:3]))
```

	x	y	z
1	1	-4	A
2	2	-5	B
3	3	-6	C


```
> my_data_frame[1:2]
```

	x	y
1	1	-4
2	2	-5
3	3	-6


```
> my_data_frame[1, 3]
```

```
[1] A
```

```
Levels: A B C
```



- If you subset using a single vector, the result behaves as a list
- If you subset using two vectors, the results behaves as a matrix
- Try the following:
 - 1 If you didn't already, set up the dataframe: `my_data_frame <- data.frame(x = 1:3, y = -4:-6, z = LETTERS[1:3])`
 - 2 What happens if you call `my_data_frame[1]`?
 - 3 What happens if you call `my_data_frame[1,]`?
 - 4 What happens if you call `my_data_frame["x"]`?
 - 5 What happens if you call `my_data_frame[, "x"]`?
 - 6 How can you return the first and third columns of the dataframe?
 - 7 What type of structure is returned when you index as a list to return one column?
 - 8 What type of structure is returned when you index as a matrix to return one column?



- You can use `colnames(df)` to access the column names of dataframe `df` (the variables, generally)
- You can use `row.names(df)` to access the row names of dataframe `df` (the observations, generally)
 - Note, tibbles discourage special row names
- You can use `dim()` to find the number of columns and rows in a dataframe (be wary of `length()`)
- `str()` is infinitely useful for understanding your dataframe



\$ is an operator that can access a variable within a dataframe
e.g., For the iris dataset, `iris$Sepal.Length` will return all sepal length values and `iris$Sepal.Length[1:3]` will return the first three sepal length values.



- Another quirk of "base" dataframes: partial column matching
 - Try it, typing `iris$Spec` returns the full column `iris$Species`
- More confusing, if partial matching would return two columns a value of "NULL" is produced
 - Try it, typing `iris$Sepal` returns `NULL`
- Tibbles do NOT have partial matching
 - First load the tibble library with `library(tibble)`
 - Now save the iris dataset as a tibble `iris2 <- as_tibble(iris)`
 - Try the above partial matches with `iris2` what happens?



- Today we will continue with tibbles and tibbles vs. data frames
- Note: New due date for HW 5 (since we will cover strings/factors Thursday)
- Review of some new resources
 - Lecture notes are available (ymmv)
 - Shared google doc



- Data frames are popular because they hold data in a structured format that makes machine learning and analysis easy
- Tibbles are a "species" of data frames that fixes some of the more error prone quirks of data frames
- We learned that we can subset (generic) data frames and tibbles using `$`, `[,]`, `[]`, and `[[]]`



- Sometimes you might want to save a column name as a variable (e.g., `sl <- "Sepal.Length"`)
- Can be useful for long or hard to type column names, or if you need to reference the column many places
- `$` only works with the column's "real" name, so if you use a variable for a column name you will need to reference with `[]`
 - Try the above, set the variable `sl <- "Sepal.Length"`
 - What happens if you reference by `iris$sl`?
 - What about `iris[, sl]`?



- We have seen subsetting with `[]` and `$`, but there is also `[[]]`
- Recall, when `[]` is used to subset a list, it returns a list, whereas when `[[]]` is used to subset a list, it returns its contents.
- Similarly, you can use `[[]]` to extract the contents of a particular column from a dataframe
- You can use column names `iris[["Species"]]` or ordinal locations `iris[[1]]` in the double square brackets
- Note: for both tibbles and dataframes `[[]]` subsetting always returns the contents of the column, not another dataframe
- Note: when using `[[]]` you can only select one (not multiple) columns



- \$ does partial matching in data frames (not tibbles)
- [[]] never does partial matching



- **Simplifying** subsets returns the simplest possible data structure that can represent the output
- **Preserving** subsets keep the output structure the same as the input
- The **drop** option when subsetting is one of the most common sources of programming error
- By default, a data frame (not a tibble) will **simplify** when returning a single column, but **preserve** when returning a single row
- You can set 'drop = F' to **preserve** even when returning a column

```
tester[1, ]  
class(tester[1, ])  
tester[,1]  
class(tester[,1])  
tester[,1, drop = F]  
class(tester[, 1, drop = F])
```




- Data frames will not allow you to use special characters like spaces or punctuation in column names
- Tibbles let you name columns anything, though if they contain special characters you must surround them with back ticks



- Tribbles also allow you to input rows using the `tribble()` function
- This can be particularly useful when adding new observations to a tibble



- Tibbles also give you more control when viewing
- Note traditional data frames will print the whole data frame, regardless of size
- Tibbles will by default print 10 rows and however many columns fit on your screen
- But this can be adjusted

Summary: Tibbles versus Data Frames



- Data frames are one of the (the?) most popular and useful data structures in R
- Tibbles aim to fix some of the finicky features of data frames, specifically
 - Tibbles will default to keeping strings as strings (data frames convert strings to factors by default)
 - Tibbles preserve all subsets made with `[]` by default (data frames simplify single column subsets by default)
 - Note, both will simplify when subsetting with `$`
 - Tibbles discourage the use of row names
 - Tibbles have a pretty output by default

Activity: Comparing Data Frames to Tibbles



Go to <https://github.com/abbiepopa/bsds100> and select **Tibbles versus Data Frames** under **Class Code**